


图片集数据增强和神经网络 train 模型

目标：飞机发动机在装机前及一定周期的飞行后需要对叶片表面是否发生损伤进行孔探检测，目前主要是靠检验人员人眼识别发动机叶片表面是否存在划痕、裂纹等异常，检测效率不高，同时受人主观判断影响较大。基于次工程问题提出通过神经网络和机器视觉的方法，对发动机叶片损伤进行智能化自动检测。

由于能获取到的各类发动机叶片损伤的数据有限，通过数据集增强来提高数据集中各类异常的样本数量，然后在神经网络训练模型中提高模型的检测精度。

图片集数据增强算法：包括水平翻转， ± 20 度随机旋转，随机调整颜色和增加锐化效果四类增强算法。


数据增强前：










数据增强后:














数据增强代码：

```
import torch
from torchvision import transforms
from PIL import Image
import os

# 定义数据增强变换
transforms_list = [
    transforms.RandomHorizontalFlip(), # 随机水平翻转
    transforms.RandomRotation(20), # 随机旋转
    transforms.ColorJitter(brightness=0.5, contrast=0.5, saturation=0.5, hue=0.5), # 随机调整
    颜色
    transforms.RandomAffine(degrees=0, shear=10, scale=(0.9, 1.1), fillcolor=0), # 锐化效果
    # 如果你还有其他变换，可以继续添加
]

# 图片所在的目录
input_dir = 'D:/bata5-5/windows/windows2/data1' # 替换为你的图片目录
output_dir = 'D:/bata5-5/windows/windows2/data1out' # 替换为保存增强图片的目录

# 确保输出目录存在
if not os.path.exists(output_dir):
    os.makedirs(output_dir)

# 遍历目录中的所有图片
for filename in os.listdir(input_dir):
    if filename.lower().endswith(('.png', '.jpg', '.jpeg', '.tiff', '.bmp', '.gif')):
```

```

image_path = os.path.join(input_dir, filename)
image = Image.open(image_path)

# 应用每个变换并保存结果
for transform in transforms_list:
    # 应用变换
    augmented_image = transform(image)
    # 保存增强后的图片
    output_path = os.path.join(output_dir,
f'{filename}_{transform.__class__.__name__}.jpg')
    augmented_image.save(output_path)

print(f'All images have been augmented and saved to {output_dir}')

```

神经网络模型训练代码：

```

import argparse
import logging
import math
import os
import random
import time
from copy import deepcopy
from pathlib import Path
from threading import Thread
import os

os.environ['KMP_DUPLICATE_LIB_OK'] = 'True'

import numpy as np
import torch.distributed as dist
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
import torch.optim.lr_scheduler as lr_scheduler
import torch.utils.data
import yaml
from torch.cuda import amp
from torch.nn.parallel import DistributedDataParallel as DDP
from torch.utils.tensorboard import SummaryWriter
from tqdm import tqdm

import test # import test.py to get mAP after each epoch
from models.experimental import attempt_load
from models.yolo import Model

```



```

from utils.autoanchor import check_anchors
from utils.datasets import create_dataloader
from utils.general import labels_to_class_weights, increment_path, labels_to_image_weights,
init_seeds, \
    fitness, strip_optimizer, get_latest_run, check_dataset, check_file, check_git_status,
check_img_size, \
    check_requirements, print_mutation, set_logging, one_cycle, colorstr
from utils.google_utils import attempt_download
from utils.loss import ComputeLoss
from utils.plots import plot_images, plot_labels, plot_results, plot_evolution
from utils.torch_utils import ModelEMA, select_device, intersect_dicts,
torch_distributed_zero_first, is_parallel
from utils.wandb_logging.wandb_utils import WandbLogger, check_wandb_resume

logger = logging.getLogger(__name__)

def train(hyp, opt, device, tb_writer=None):
    logger.info(colorstr('hyperparameters: ') + ', '.join(f'{k}={v}' for k, v in hyp.items()))
    save_dir, epochs, batch_size, total_batch_size, weights, rank = \
        Path(opt.save_dir), opt.epochs, opt.batch_size, opt.total_batch_size, opt.weights,
opt.global_rank

    # Directories
    wdir = save_dir / 'weights'
    wdir.mkdir(parents=True, exist_ok=True) # make dir
    last = wdir / 'last.pt'
    best = wdir / 'best.pt'
    results_file = save_dir / 'results.txt'

    # Save run settings
    with open(save_dir / 'hyp.yaml', 'w') as f:
        yaml.dump(hyp, f, sort_keys=False)
    with open(save_dir / 'opt.yaml', 'w') as f:
        yaml.dump(vars(opt), f, sort_keys=False)

    # Configure
    plots = not opt.evolve # create plots
    cuda = device.type != 'cpu'
    init_seeds(2 + rank)
    with open(opt.data) as f:
        data_dict = yaml.load(f, Loader=yaml.SafeLoader) # data dict
    is_coco = opt.data.endswith('coco.yaml')

```

```

# Logging- Doing this before checking the dataset. Might update data_dict
loggers = {'wandb': None} # loggers dict
if rank in [-1, 0]:
    opt.hyp = hyp # add hyperparameters
    run_id = torch.load(weights).get('wandb_id') if weights.endswith('.pt') and
os.path.isfile(weights) else None
    wandb_logger = WandbLogger(opt, Path(opt.save_dir).stem, run_id, data_dict)
    loggers['wandb'] = wandb_logger.wandb
    data_dict = wandb_logger.data_dict
    if wandb_logger.wandb:
        weights, epochs, hyp = opt.weights, opt.epochs, opt.hyp # WandbLogger might
update weights, epochs if resuming

    nc = 1 if opt.single_cls else int(data_dict['nc']) # number of classes
    names = ['item'] if opt.single_cls and len(data_dict['names']) != 1 else data_dict['names'] #
class names
    assert len(names) == nc, '%g names found for nc=%g dataset in %s' % (len(names), nc,
opt.data) # check

# Model
pretrained = weights.endswith('.pt')
if pretrained:
    with torch_distributed_zero_first(rank):
        attempt_download(weights) # download if not found locally
        ckpt = torch.load(weights, map_location=device) # load checkpoint
        model = Model(opt.cfg or ckpt['model'].yaml, ch=3, nc=nc,
anchors=hyp.get('anchors')).to(device) # create
        exclude = ['anchor'] if (opt.cfg or hyp.get('anchors')) and not opt.resume else [] #
exclude keys
        state_dict = ckpt['model'].float().state_dict() # to FP32
        state_dict = intersect_dicts(state_dict, model.state_dict(), exclude=exclude) # intersect
        model.load_state_dict(state_dict, strict=False) # load
        logger.info("Transferred %g/%g items from %s' % (len(state_dict),
len(model.state_dict()), weights)) # report
    else:
        model = Model(opt.cfg, ch=3, nc=nc, anchors=hyp.get('anchors')).to(device) # create
        with torch_distributed_zero_first(rank):
            check_dataset(data_dict) # check
            train_path = data_dict['train']
            test_path = data_dict['val']

# Freeze
freeze = [] # parameter names to freeze (full or partial)
for k, v in model.named_parameters():

```

```

        v.requires_grad = True # train all layers
    if any(x in k for x in freeze):
        print('freezing %s' % k)
        v.requires_grad = False

# Optimizer
nbs = 64 # nominal batch size
accumulate = max(round(nbs / total_batch_size), 1) # accumulate loss before optimizing
hyp['weight_decay'] *= total_batch_size * accumulate / nbs # scale weight_decay
logger.info(f"Scaled weight_decay = {hyp['weight_decay']}")

pg0, pg1, pg2 = [], [], [] # optimizer parameter groups
for k, v in model.named_modules():
    if hasattr(v, 'bias') and isinstance(v.bias, nn.Parameter):
        pg2.append(v.bias) # biases
    if isinstance(v, nn.BatchNorm2d):
        pg0.append(v.weight) # no decay
    elif hasattr(v, 'weight') and isinstance(v.weight, nn.Parameter):
        pg1.append(v.weight) # apply decay

if opt.adam:
    optimizer = optim.Adam(pg0, lr=hyp['lr0'], betas=(hyp['momentum'], 0.999)) # adjust
    # adjust beta1 to momentum
    else:
        optimizer = optim.SGD(pg0, lr=hyp['lr0'], momentum=hyp['momentum'],
                                nesterov=True)

    optimizer.add_param_group({'params': pg1, 'weight_decay': hyp['weight_decay']}) # add
    # add pg1 with weight_decay
    optimizer.add_param_group({'params': pg2}) # add pg2 (biases)
    logger.info('Optimizer groups: %g .bias, %g conv.weight, %g other' % (len(pg2), len(pg1),
    len(pg0)))
    del pg0, pg1, pg2

# Scheduler https://arxiv.org/pdf/1812.01187.pdf
# https://pytorch.org/docs/stable/_modules/torch/optim/lr_scheduler.html#OneCycleLR
if opt.linear_lr:
    lf = lambda x: (1 - x / (epochs - 1)) * (1.0 - hyp['lrf']) + hyp['lrf'] # linear
    else:
        lf = one_cycle(1, hyp['lrf'], epochs) # cosine 1->hyp['lrf']
    scheduler = lr_scheduler.LambdaLR(optimizer, lr_lambda=lf)
    # plot_lr_scheduler(optimizer, scheduler, epochs)

# EMA

```

```

ema = ModelEMA(model) if rank in [-1, 0] else None

# Resume
start_epoch, best_fitness = 0, 0.0
if pretrained:
    # Optimizer
    if ckpt['optimizer'] is not None:
        optimizer.load_state_dict(ckpt['optimizer'])
        best_fitness = ckpt['best_fitness']

    # EMA
    if ema and ckpt.get('ema'):
        ema.ema.load_state_dict(ckpt['ema'].float().state_dict())
        ema.updates = ckpt['updates']

    # Results
    if ckpt.get('training_results') is not None:
        results_file.write_text(ckpt['training_results'])  # write results.txt

    # Epochs
    start_epoch = ckpt['epoch'] + 1
    if opt.resume:
        assert start_epoch > 0, '%s training to %g epochs is finished, nothing to resume.' %
(weights, epochs)
    if epochs < start_epoch:
        logger.info('%s has been trained for %g epochs. Fine-tuning for %g additional
epochs.' %
                    (weights, ckpt['epoch'], epochs))
        epochs += ckpt['epoch']  # finetune additional epochs

    del ckpt, state_dict

# Image sizes
gs = max(int(model.stride.max()), 32)  # grid size (max stride)
nl = model.model[-1].nl  # number of detection layers (used for scaling hyp['obj'])
imgsz, imgsz_test = [check_img_size(x, gs) for x in opt.img_size]  # verify imgsz are
gs-multiples

# DP mode
if cuda and rank == -1 and torch.cuda.device_count() > 1:
    model = torch.nn.DataParallel(model)

# SyncBatchNorm
if opt.sync_bn and cuda and rank != -1:

```

```

model = torch.nn.SyncBatchNorm.convert_sync_batchnorm(model).to(device)
logger.info('Using SyncBatchNorm()')

# Trainloader
dataloader, dataset = create_dataloader(train_path, imgsz, batch_size, gs, opt,
                                       hyp=hyp,                        augment=True,
cache=opt.cache_images, rect=opt.rect, rank=rank,
                                       world_size=opt.world_size,
workers=opt.workers,
                                       image_weights=opt.image_weights,
quad=opt.quad, prefix=colorstr('train: '))
mlc = np.concatenate(dataset.labels, 0)[: , 0].max() # max label class
nb = len(dataloader) # number of batches
assert mlc < nc, 'Label class %g exceeds nc=%g in %s. Possible class labels are 0-%g' %
(mlc, nc, opt.data, nc - 1)

# Process 0
if rank in [-1, 0]:
    testloader = create_dataloader(test_path, imgsz_test, batch_size * 2, gs, opt, #
testloader
                                hyp=hyp, cache=opt.cache_images and not
opt.notest, rect=True, rank=-1,
                                world_size=opt.world_size,
workers=opt.workers,
                                pad=0.5, prefix=colorstr('val: '))[0]

if not opt.resume:
    labels = np.concatenate(dataset.labels, 0)
    c = torch.tensor(labels[:, 0]) # classes
    # cf = torch.bincount(c.long(), minlength=nc) + 1. # frequency
    # model._initialize_biases(cf.to(device))
    if plots:
        plot_labels(labels, names, save_dir, loggers)
        if tb_writer:
            tb_writer.add_histogram('classes', c, 0)

# Anchors
if not opt.noautoanchor:
    check_anchors(dataset, model=model, thr=hyp['anchor_t'], imgsz=imgsz)
    model.half().float() # pre-reduce anchor precision

# DDP mode
if cuda and rank != -1:
    model = DDP(model, device_ids=[opt.local_rank], output_device=opt.local_rank,

```

```

# nn.MultiheadAttention incompatibility with DDP
https://github.com/pytorch/pytorch/issues/26698
find_unused_parameters=any(isinstance(layer, nn.MultiheadAttention)
for layer in model.modules()))

```

```

# Model parameters
hyp['box'] *= 3. / nl # scale to layers
hyp['cls'] *= nc / 80. * 3. / nl # scale to classes and layers
hyp['obj'] *= (imgsz / 640) ** 2 * 3. / nl # scale to image size and layers
hyp['label_smoothing'] = opt.label_smoothing
model.nc = nc # attach number of classes to model
model.hyp = hyp # attach hyperparameters to model
model.gr = 1.0 # iou loss ratio (obj_loss = 1.0 or iou)
model.class_weights = labels_to_class_weights(dataset.labels, nc).to(device) * nc # attach
class weights
model.names = names

```

```

# Start training
t0 = time.time()
nw = max(round(hyp['warmup_epochs'] * nb), 1000) # number of warmup iterations,
max(3 epochs, 1k iterations)
# nw = min(nw, (epochs - start_epoch) / 2 * nb) # limit warmup to < 1/2 of training
maps = np.zeros(nc) # mAP per class
results = (0, 0, 0, 0, 0, 0, 0) # P, R, mAP@.5, mAP@.5-.95, val_loss(box, obj, cls)
scheduler.last_epoch = start_epoch - 1 # do not move
scaler = amp.GradScaler(enabled=cuda)
compute_loss = ComputeLoss(model) # init loss class
logger.info(f'Image sizes {imgsz} train, {imgsz_test} test\n'
            f'Using {dataloader.num_workers} dataloader workers\n'
            f'Logging results to {save_dir}\n'
            f'Starting training for {epochs} epochs...')
for epoch in range(start_epoch, epochs): # epoch

```

```

-----
    model.train()

```

```

    # Update image weights (optional)
    if opt.image_weights:
        # Generate indices
        if rank in [-1, 0]:
            cw = model.class_weights.cpu().numpy() * (1 - maps) ** 2 / nc # class
weights
            iw = labels_to_image_weights(dataset.labels, nc=nc, class_weights=cw) #
image weights
            dataset.indices = random.choices(range(dataset.n), weights=iw, k=dataset.n)

```

```

# rand weighted idx
    # Broadcast if DDP
    if rank != -1:
        indices = (torch.tensor(dataset.indices) if rank == 0 else
torch.zeros(dataset.n)).int()
        dist.broadcast(indices, 0)
    if rank != 0:
        dataset.indices = indices.cpu().numpy()

# Update mosaic border
# b = int(random.uniform(0.25 * imgsz, 0.75 * imgsz + gs) // gs * gs)
# dataset.mosaic_border = [b - imgsz, -b] # height, width borders

mloss = torch.zeros(4, device=device) # mean losses
if rank != -1:
    dataloader.sampler.set_epoch(epoch)
    pbar = enumerate(dataloader)
    logger.info(("n" + "%10s" * 8) % ('Epoch', 'gpu_mem', 'box', 'obj', 'cls', 'total', 'labels',
'img_size'))
    if rank in [-1, 0]:
        pbar = tqdm(pbar, total=nb) # progress bar
        optimizer.zero_grad()
        for i, (imgs, targets, paths, _) in pbar: # batch
            -----
            ni = i + nb * epoch # number integrated batches (since train start)
            imgs = imgs.to(device, non_blocking=True).float() / 255.0 # uint8 to float32,
0-255 to 0.0-1.0

            # Warmup
            if ni <= nw:
                xi = [0, nw] # x interp
                # model.gr = np.interp(ni, xi, [0.0, 1.0]) # iou loss ratio (obj_loss = 1.0 or
iou)

                accumulate = max(1, np.interp(ni, xi, [1, nbs / total_batch_size])).round()
                for j, x in enumerate(optimizer.param_groups):
                    # bias lr falls from 0.1 to lr0, all other lrs rise from 0.0 to lr0
                    x['lr'] = np.interp(ni, xi, [hyp['warmup_bias_lr'] if j == 2 else 0.0,
x['initial_lr'] * lf(epoch)])
                    if 'momentum' in x:
                        x['momentum'] = np.interp(ni, xi, [hyp['warmup_momentum'],
hyp['momentum']]))

            # Multi-scale
            if opt.multi_scale:

```



```

sz = random.randrange(imgsz * 0.5, imgsz * 1.5 + gs) // gs * gs  # size
sf = sz / max(imgs.shape[2:])  # scale factor
if sf != 1:
    ns = [math.ceil(x * sf / gs) * gs for x in imgs.shape[2:]]  # new shape
(stretched to gs-multiple)
    imgs = F.interpolate(imgs, size=ns, mode='bilinear', align_corners=False)

# Forward
with amp.autocast(enabled=cuda):
    pred = model(imgs)  # forward
    loss, loss_items = compute_loss(pred, targets.to(device))  # loss scaled by
batch_size

    if rank != -1:
        loss *= opt.world_size  # gradient averaged between devices in DDP
mode

    if opt.quad:
        loss *= 4.

# Backward
scaler.scale(loss).backward()

# Optimize
if ni % accumulate == 0:
    scaler.step(optimizer)  # optimizer.step
    scaler.update()
    optimizer.zero_grad()
    if ema:
        ema.update(model)

# Print
if rank in [-1, 0]:
    mloss = (mloss * i + loss_items) / (i + 1)  # update mean losses
    mem = '%.3gG' % (torch.cuda.memory_reserved() / 1E9 if
torch.cuda.is_available() else 0)  # (GB)
    s = ('%10s' * 2 + '%10.4g' * 6) % (
        '%g/%g' % (epoch, epochs - 1), mem, *mloss, targets.shape[0],
imgs.shape[-1])
    pbar.set_description(s)

# Plot
if plots and ni < 3:
    f = save_dir / f'train_batch{ni}.jpg'  # filename
    Thread(target=plot_images, args=(imgs, targets, paths, f),
daemon=True).start()

```

```

        # if tb_writer:
        #             tb_writer.add_image(f, result, dataformats='HWC',
global_step=epoch)
        #             tb_writer.add_graph(torch.jit.trace(model, imgs, strict=False), [])
# add model graph
        elif plots and ni == 10 and wandb_logger.wandb:
            wandb_logger.log({"Mosaics": [wandb_logger.wandb.Image(str(x),
caption=x.name) for x in
                                save_dir.glob('train*.jpg') if
x.exists()]})

        #                                     end                                batch
-----
        #                                     end                                epoch
-----

# Scheduler
lr = [x['lr'] for x in optimizer.param_groups] # for tensorboard
scheduler.step()

# DDP process 0 or single-GPU
if rank in [-1, 0]:
    # mAP
    ema.update_attr(model, include=['yaml', 'nc', 'hyp', 'gr', 'names', 'stride',
'class_weights'])
    final_epoch = epoch + 1 == epochs
    if not opt.notest or final_epoch: # Calculate mAP
        wandb_logger.current_epoch = epoch + 1
        results, maps, times = test.test(data_dict,
                                         batch_size=batch_size * 2,
                                         imsz=imsz_test,
                                         model=ema.ema,
                                         single_cls=opt.single_cls,
                                         dataloader=testloader,
                                         save_dir=save_dir,
                                         verbose=nc < 50 and final_epoch,
                                         plots=plots and final_epoch,
                                         wandb_logger=wandb_logger,
                                         compute_loss=compute_loss,
                                         is_coco=is_coco)

# Write
with open(results_file, 'a') as f:
    f.write(s + '%10.4g' * 7 % results + '\n') # append metrics, val_loss

```

```

        if len(opt.name) and opt.bucket:
            os.system('gsutil cp %s gs://%s/results/results%s.txt' % (results_file,
opt.bucket, opt.name))

    # Log
    tags = ['train/box_loss', 'train/obj_loss', 'train/cls_loss', # train loss
            'metrics/precision', 'metrics/recall', 'metrics/mAP_0.5',
'metrics/mAP_0.5:0.95',
            'val/box_loss', 'val/obj_loss', 'val/cls_loss', # val loss
            'x/lr0', 'x/lr1', 'x/lr2'] # params
    for x, tag in zip(list(mloss[:-1]) + list(results) + lr, tags):
        if tb_writer:
            tb_writer.add_scalar(tag, x, epoch) # tensorboard
        if wandb_logger.wandb:
            wandb_logger.log({tag: x}) # W&B

    # Update best mAP
    fi = fitness(np.array(results).reshape(1, -1)) # weighted combination of [P, R,
mAP@.5, mAP@.5-.95]
    if fi > best_fitness:
        best_fitness = fi
    wandb_logger.end_epoch(best_result=best_fitness == fi)

    # Save model
    if (not opt.nosave) or (final_epoch and not opt.evolve): # if save
        ckpt = {'epoch': epoch,
                'best_fitness': best_fitness,
                'training_results': results_file.read_text(),
                'model': deepcopy(model.module if is_parallel(model) else
model).half(),
                'ema': deepcopy(ema.ema).half(),
                'updates': ema.updates,
                'optimizer': optimizer.state_dict(),
                'wandb_id': wandb_logger.wandb_run.id if wandb_logger.wandb
else None}

        # Save last, best and delete
        torch.save(ckpt, last)
        if best_fitness == fi:
            torch.save(ckpt, best)
        if wandb_logger.wandb:
            if ((epoch + 1) % opt.save_period == 0 and not final_epoch) and
opt.save_period != -1:
                wandb_logger.log_model(

```

```

                                last.parent, opt, epoch, fi, best_model=best_fitness == fi)
                                del ckpt

                                #                                end                                epoch
                                -----
                                # end training
                                if rank in [-1, 0]:
                                    # Plots
                                    if plots:
                                        plot_results(save_dir=save_dir) # save as results.png
                                        if wandb_logger.wandb:
                                            files = ['results.png', 'confusion_matrix.png', *[f'{x}_curve.png' for x in ('F1',
'PR', 'P', 'R')]]
                                            wandb_logger.log({"Results": [wandb_logger.wandb.Image(str(save_dir / f),
caption=f) for f in files
                                                            if (save_dir / f).exists()]}))

                                    # Test best.pt
                                    logger.info("%g epochs completed in %.3f hours.\n" % (epoch - start_epoch + 1,
(time.time() - t0) / 3600))
                                    if opt.data.endswith('coco.yaml') and nc == 80: # if COCO
                                        for m in (last, best) if best.exists() else (last): # speed, mAP tests
                                            results, _, _ = test.test(opt.data,
                                                                    batch_size=batch_size * 2,
                                                                    imgsiz=imgsiz_test,
                                                                    conf_thres=0.001,
                                                                    iou_thres=0.7,
                                                                    model=attempt_load(m, device).half(),
                                                                    single_cls=opt.single_cls,
                                                                    dataloader=testloader,
                                                                    save_dir=save_dir,
                                                                    save_json=True,
                                                                    plots=False,
                                                                    is_coco=is_coco)

                                    # Strip optimizers
                                    final = best if best.exists() else last # final model
                                    for f in last, best:
                                        if f.exists():
                                            strip_optimizer(f) # strip optimizers
                                    if opt.bucket:
                                        os.system(f'gsutil cp {final} gs://{opt.bucket}/weights') # upload
                                    if wandb_logger.wandb and not opt.evolve: # Log the stripped model
                                        wandb_logger.wandb.log_artifact(str(final), type='model',
                                                                    name='run_' + wandb_logger.wandb_run.id

```

```

+ '_model',

                                                                    aliases=['last', 'best', 'stripped'])

        wandb_logger.finish_run()
    else:
        dist.destroy_process_group()
    torch.cuda.empty_cache()
    return results

if __name__ == '__main__':
    parser = argparse.ArgumentParser()
    parser.add_argument('--weights', type=str, default='weights/yolov5s.pt', help='initial weights
path')
    parser.add_argument('--cfg', type=str, default='models/blade.yaml', help='model.yaml path')
    parser.add_argument('--data', type=str, default='data/blade.yaml', help='data.yaml path')
    parser.add_argument('--hyp', type=str, default='data/hyp.scratch.yaml', help='hyperparameters
path')
    parser.add_argument('--epochs', type=int, default=1000)
    parser.add_argument('--batch-size', type=int, default=16, help='total batch size for all GPUs')
    parser.add_argument('--img-size', nargs='+', type=int, default=[640, 640], help='[train, test]
image sizes')
    parser.add_argument('--rect', action='store_true', help='rectangular training')
    parser.add_argument('--resume', nargs='?', const=True, default=False, help='resume most
recent training')
    parser.add_argument('--nosave', action='store_true', help='only save final checkpoint')
    parser.add_argument('--notest', action='store_true', help='only test final epoch')
    parser.add_argument('--noautoanchor', action='store_true', help='disable autoanchor check')
    parser.add_argument('--evolve', action='store_true', help='evolve hyperparameters')
    parser.add_argument('--bucket', type=str, default="", help='gsutil bucket')
    parser.add_argument('--cache-images', action='store_true', help='cache images for faster
training')
    parser.add_argument('--image-weights', action='store_true', help='use weighted image
selection for training')
    parser.add_argument('--device', default='0', help='cuda device, i.e. 0 or 0,1,2,3 or cpu')
    parser.add_argument('--multi-scale', action='store_true', help='vary img-size +/- 50%%')
    parser.add_argument('--single-cls', action='store_true', help='train multi-class data as
single-class')
    parser.add_argument('--adam', action='store_true', help='use torch.optim.Adam() optimizer')
    parser.add_argument('--sync-bn', action='store_true', help='use SyncBatchNorm, only
available in DDP mode')
    parser.add_argument('--local_rank', type=int, default=-1, help='DDP parameter, do not
modify')
    parser.add_argument('--workers', type=int, default=8, help='maximum number of dataloader
workers')

```

```

parser.add_argument('--project', default='runs/train', help='save to project/name')
parser.add_argument('--entity', default=None, help='W&B entity')
parser.add_argument('--name', default='exp', help='save to project/name')
parser.add_argument('--exist-ok', action='store_true', help='existing project/name ok, do not
increment')
parser.add_argument('--quad', action='store_true', help='quad dataloader')
parser.add_argument('--linear-lr', action='store_true', help='linear LR')
parser.add_argument('--label-smoothing', type=float, default=0.0, help='Label smoothing
epsilon')
parser.add_argument('--upload_dataset', action='store_true', help='Upload dataset as W&B
artifact table')
parser.add_argument('--bbox_interval', type=int, default=-1, help='Set bounding-box image
logging interval for W&B')
parser.add_argument('--save_period', type=int, default=-1, help='Log model after every
"save_period" epoch')
parser.add_argument('--artifact_alias', type=str, default="latest", help='version of dataset
artifact to be used')
opt = parser.parse_args()

# Set DDP variables
opt.world_size = int(os.environ['WORLD_SIZE']) if 'WORLD_SIZE' in os.environ else 1
opt.global_rank = int(os.environ['RANK']) if 'RANK' in os.environ else -1
set_logging(opt.global_rank)
if opt.global_rank in [-1, 0]:
    check_git_status()
    check_requirements()

# Resume
wandb_run = check_wandb_resume(opt)
if opt.resume and not wandb_run: # resume an interrupted run
    ckpt = opt.resume if isinstance(opt.resume, str) else get_latest_run() # specified or
most recent path
    assert os.path.isfile(ckpt), 'ERROR: --resume checkpoint does not exist'
    apriori = opt.global_rank, opt.local_rank
    with open(Path(ckpt).parent.parent / 'opt.yaml') as f:
        opt = argparse.Namespace(**yaml.load(f, Loader=yaml.SafeLoader)) # replace
        opt.cfg, opt.weights, opt.resume, opt.batch_size, opt.global_rank, opt.local_rank = ",
ckpt, True, opt.total_batch_size, *apriori # reinstate
        logger.info('Resuming training from %s' % ckpt)
    else:
        # opt.hyp = opt.hyp or ('hyp.finetune.yaml' if opt.weights else 'hyp.scratch.yaml')
        opt.data, opt.cfg, opt.hyp = check_file(opt.data), check_file(opt.cfg), check_file(opt.hyp)
# check files
assert len(opt.cfg) or len(opt.weights), 'either --cfg or --weights must be specified'

```

```

        opt.img_size.extend([opt.img_size[-1]] * (2 - len(opt.img_size))) # extend to 2 sizes
(train, test)
    opt.name = 'evolve' if opt.evolve else opt.name
    opt.save_dir = increment_path(Path(opt.project) / opt.name, exist_ok=opt.exist_ok |
opt.evolve) # increment run

# DDP mode
opt.total_batch_size = opt.batch_size
device = select_device(opt.device, batch_size=opt.batch_size)
if opt.local_rank != -1:
    assert torch.cuda.device_count() > opt.local_rank
    torch.cuda.set_device(opt.local_rank)
    device = torch.device('cuda', opt.local_rank)
    dist.init_process_group(backend='nccl', init_method='env://') # distributed backend
    assert opt.batch_size % opt.world_size == 0, '--batch-size must be multiple of CUDA
device count'
    opt.batch_size = opt.total_batch_size // opt.world_size

# Hyperparameters
with open(opt.hyp) as f:
    hyp = yaml.load(f, Loader=yaml.SafeLoader) # load hyps

# Train
logger.info(opt)
if not opt.evolve:
    tb_writer = None # init loggers
    if opt.global_rank in [-1, 0]:
        prefix = colorstr('tensorboard: ')
        logger.info(f'{prefix}Start with 'tensorboard --logdir {opt.project}', view at
http://localhost:6006/")
        tb_writer = SummaryWriter(opt.save_dir) # Tensorboard
    train(hyp, opt, device, tb_writer)

# Evolve hyperparameters (optional)
else:
    # Hyperparameter evolution metadata (mutation scale 0-1, lower_limit, upper_limit)
    meta = {'lr0': (1, 1e-5, 1e-1), # initial learning rate (SGD=1E-2, Adam=1E-3)
            'lrf': (1, 0.01, 1.0), # final OneCycleLR learning rate (lr0 * lrf)
            'momentum': (0.3, 0.6, 0.98), # SGD momentum/Adam beta1
            'weight_decay': (1, 0.0, 0.001), # optimizer weight decay
            'warmup_epochs': (1, 0.0, 5.0), # warmup epochs (fractions ok)
            'warmup_momentum': (1, 0.0, 0.95), # warmup initial momentum
            'warmup_bias_lr': (1, 0.0, 0.2), # warmup initial bias lr
            'box': (1, 0.02, 0.2), # box loss gain

```



```

'cls': (1, 0.2, 4.0), # cls loss gain
'cls_pw': (1, 0.5, 2.0), # cls BCELoss positive_weight
'obj': (1, 0.2, 4.0), # obj loss gain (scale with pixels)
'obj_pw': (1, 0.5, 2.0), # obj BCELoss positive_weight
'iou_t': (0, 0.1, 0.7), # IoU training threshold
'anchor_t': (1, 2.0, 8.0), # anchor-multiple threshold
'anchors': (2, 2.0, 10.0), # anchors per output grid (0 to ignore)
'fl_gamma': (0, 0.0, 2.0), # focal loss gamma (efficientDet default
gamma=1.5)
'hsv_h': (1, 0.0, 0.1), # image HSV-Hue augmentation (fraction)
'hsv_s': (1, 0.0, 0.9), # image HSV-Saturation augmentation (fraction)
'hsv_v': (1, 0.0, 0.9), # image HSV-Value augmentation (fraction)
'degrees': (1, 0.0, 45.0), # image rotation (+/- deg)
'translate': (1, 0.0, 0.9), # image translation (+/- fraction)
'scale': (1, 0.0, 0.9), # image scale (+/- gain)
'shear': (1, 0.0, 10.0), # image shear (+/- deg)
'perspective': (0, 0.0, 0.001), # image perspective (+/- fraction), range
0-0.001
'flipud': (1, 0.0, 1.0), # image flip up-down (probability)
'fliplr': (0, 0.0, 1.0), # image flip left-right (probability)
'mosaic': (1, 0.0, 1.0), # image mixup (probability)
'mixup': (1, 0.0, 1.0)} # image mixup (probability)

assert opt.local_rank == -1, 'DDP mode not implemented for --evolve'
opt.notest, opt.nosave = True, True # only test/save final epoch
# ei = [isinstance(x, (int, float)) for x in hyp.values()] # evolvable indices
yaml_file = Path(opt.save_dir) / 'hyp_evolved.yaml' # save best result here
if opt.bucket:
    os.system('gsutil cp gs://%s/evolve.txt .' % opt.bucket) # download evolve.txt if
exists

for _ in range(300): # generations to evolve
    if Path('evolve.txt').exists(): # if evolve.txt exists: select best hyps and mutate
        # Select parent(s)
        parent = 'single' # parent selection method: 'single' or 'weighted'
        x = np.loadtxt('evolve.txt', ndmin=2)
        n = min(5, len(x)) # number of previous results to consider
        x = x[np.argsort(-fitness(x))][:n] # top n mutations
        w = fitness(x) - fitness(x).min() # weights
        if parent == 'single' or len(x) == 1:
            # x = x[random.randint(0, n - 1)] # random selection
            x = x[random.choices(range(n), weights=w)[0]] # weighted selection
        elif parent == 'weighted':
            x = (x * w.reshape(n, 1)).sum(0) / w.sum() # weighted combination

```

```

# Mutate
mp, s = 0.8, 0.2 # mutation probability, sigma
npr = np.random
npr.seed(int(time.time()))
g = np.array([x[0] for x in meta.values()]) # gains 0-1
ng = len(meta)
v = np.ones(ng)
while all(v == 1): # mutate until a change occurs (prevent duplicates)
    v = (g * (npr.random(ng) < mp) * npr.randn(ng) * npr.random() * s +
1).clip(0.3, 3.0)

    for i, k in enumerate(hyp.keys()): # plt.hist(v.ravel(), 300)
        hyp[k] = float(x[i + 7] * v[i]) # mutate

# Constrain to limits
for k, v in meta.items():
    hyp[k] = max(hyp[k], v[1]) # lower limit
    hyp[k] = min(hyp[k], v[2]) # upper limit
    hyp[k] = round(hyp[k], 5) # significant digits

# Train mutation
results = train(hyp.copy(), opt, device)

# Write mutation results
print_mutation(hyp.copy(), results, yaml_file, opt.bucket)

# Plot results
plot_evolution(yaml_file)
print(f'Hyperparameter evolution complete. Best results saved as: {yaml_file}\n'
      f'Command to train a new model with these hyperparameters: $ python train.py
--hyp {yaml_file}')

```

检测效果:



